

Assignment - 3

Python Programming

Assignment Date	4 October 2022
Student Name	Prateeksha PK
Student Roll Number	CITC1907031
Maximum Marks	2 Marks

Problem Statement: Abalone Age Prediction

Importing Modules

```
In [ ]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
```

1. Dataset Download

```
In [ ]: #Name of the dataset: abalone.csv
```

2. Loading the Dataset into the tool

```
In [ ]: data=pd.read_csv("abalone.csv")
data.head()
```

```
Out[ ]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

Shape of data:

```
In [ ]: data.shape
```

```
Out[ ]: (4177, 9)
```

Adding 'Age' into 'Rings' - adding '1.5' to the ring data

```
In [ ]: Age=1.5+data.Rings
data["Age"]=Age
data=data.rename(columns = {'Whole weight': 'Whole_weight', 'Shucked weight': 'Shucked_weight', 'Viscera weight': 'Viscera_weight',
                             'Shell weight': 'Shell_weight'})
data=data.drop(columns=["Rings"],axis=1)
data.head()
```

```
Out[ ]:
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5


3. Perform below visualizations:

(i)Univariate Analysis

Histogram


```
In [ ]: sns.displot(data["Age"], color='darkorange')
```

```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x7fd3f837a430>
```




```
In [ ]: sns.histplot(y=data.Age,color='pink')
```

```
Out[ ]: <AxesSubplot:xlabel='Count', ylabel='Age'>
```



```
In [ ]: sns.histplot(x=data.Age,color='green')
```


```
Out[ ]: <AxesSubplot:xlabel='Age', ylabel='Count'>
```



Boxplot

```
In [ ]: sns.boxplot(x=data.Age,color='pink')
```


```
Out[ ]: <AxesSubplot:xlabel='Age'>
```



Countplot

```
In [ ]: sns.countplot(x=data.Age)
```

```
Out[ ]: <AxesSubplot:xlabel='Age', ylabel='count'>
```




(ii) Bi-Variate Analysis

Barplot

```
In [ ]: sns.barplot(x=data.Height,y=data.Age)
```


```
Out[ ]: <AxesSubplot:xlabel='Height', ylabel='Age'>
```



Linearplot

```
In [ ]: sns.lineplot(x=data.Age,y=data.Height, color='purple')
```


```
Out[ ]: <AxesSubplot:xlabel='Age', ylabel='Height'>
```



Scatterplot

```
In [ ]: sns.scatterplot(x=data.Age,y=data.Height,color='green')
```


```
Out[ ]: <AxesSubplot:xlabel='Age', ylabel='Height'>
```



Pointplot

```
In [ ]: sns.pointplot(x=data.Age, y=data.Height, color="blue")
```


```
Out[ ]: <AxesSubplot:xlabel='Age', ylabel='Height'>
```



Regplot

```
In [ ]: sns.regplot(x=data.Age,y=data.Height,color='orange')
```

```
Out[ ]: <AxesSubplot:xlabel='Age', ylabel='Height'>
```

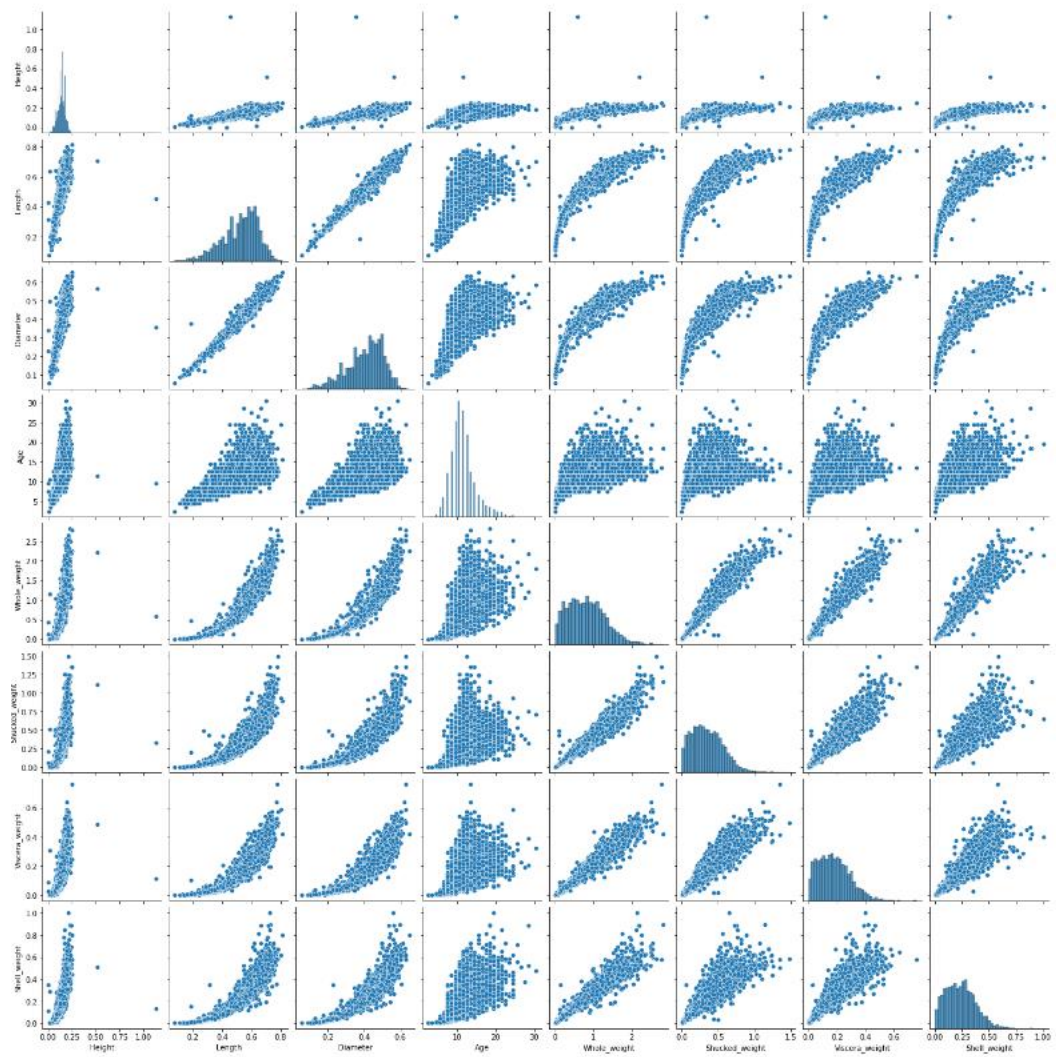


(iii) Multi-Variate Analysis

Pairplot

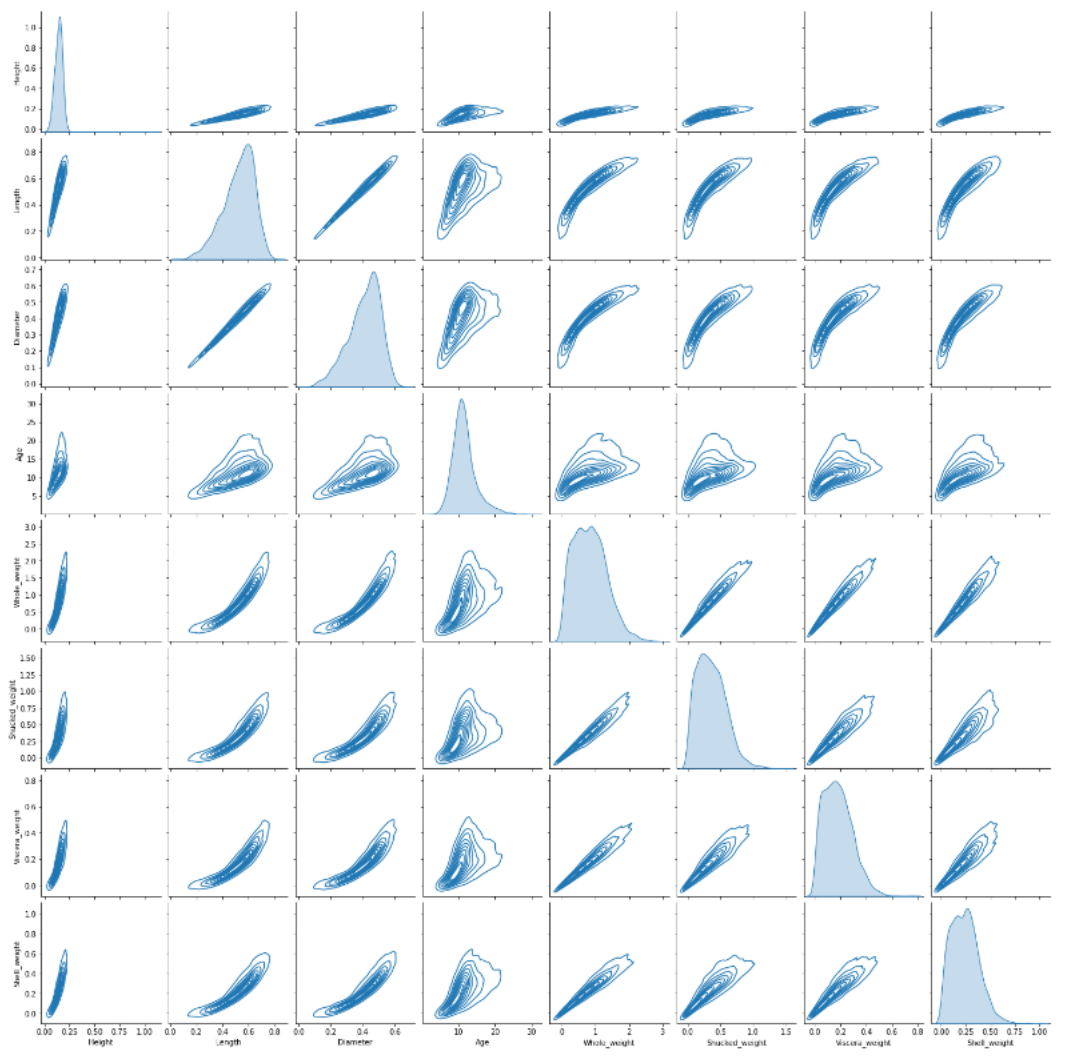
```
In [ ]: sns.pairplot(data=data[["Height","Length","Diameter","Age","Whole_weight","Shucked_weight","Viscera_weight","Shell_weight"]])
```

```
Out[ ]: <seaborn.axisgrid.PairGrid at 0x7fd3d93e1040>
```



```
In [ ]: sns.pairplot(data=data[["Height", "Length", "Diameter", "Age", "Whole_weight", "Shucked_weight", "Viscera_weight", "Shell_weight"]], kind="kde")

Out[ ]: <seaborn.axisgrid.PairGrid at 0x7fd39840c790>
```



4. Perform descriptive statistics on the dataset

```
In [ ]: data.describe(include='all')
```

Out []:	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
count	4177	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
unique	3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
top	M	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
freq	1528	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
mean	NaN	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	11.433684
std	NaN	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
min	NaN	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	2.500000
25%	NaN	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	9.500000
50%	NaN	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	10.500000
75%	NaN	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	12.500000
max	NaN	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	30.500000

5. Check for Missing values and deal with them


```
In [ ]: data.isnull().sum()
```

```
Out[ ]: Sex          0
Length        0
Diameter      0
Height        0
Whole_weight  0
Shucked_weight 0
Viscera_weight 0
Shell_weight  0
Age           0
dtype: int64
```

6. Find the outliers and replace them outliers

```
In [ ]: outliers=data.quantile(q=(0.25,0.75))
outliers
```


```
Out[ ]:      Length  Diameter  Height  Whole_weight  Shucked_weight  Viscera_weight  Shell_weight  Age
0.25    0.450      0.35    0.115      0.4415      0.186      0.0935      0.130  9.5
0.75    0.615      0.48    0.165      1.1530      0.502      0.2530      0.329  12.5
```

```
In [ ]: a = data.Age.quantile(0.25)
b = data.Age.quantile(0.75)
c = b - a
lower_limit = a - 1.5 * c
data.median(numeric_only=True)
```

```
Out[ ]: Length          0.5450
Diameter              0.4250
Height                0.1400
Whole_weight          0.7995
Shucked_weight        0.3360
Viscera_weight         0.1710
Shell_weight           0.2340
Age                  10.5000
dtype: float64
```

```
In [ ]: data['Age'] = np.where(data['Age'] < lower_limit, 7, data['Age'])
sns.boxplot(x=data.Age,showfliers = False)
```

```
Out[ ]: <AxesSubplot:xlabel='Age'>
```



7. Check for Categorical columns and perform encoding

```
In [ ]: data.head()
```

```
Out[ ]:
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

```
In [ ]: from sklearn.preprocessing import LabelEncoder

lab = LabelEncoder()
data.Sex = lab.fit_transform(data.Sex)

data.head()
```

```
Out[ ]:
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	12
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	4
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	6
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	7
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	4

8. Split the data into dependent and independent variables

```
In [ ]: y = data["Sex"]
y.head()
```

```
Out[ ]:
```

0	2
1	2
2	0
3	2
4	1

Name: Sex, dtype: int64

```
In [ ]: x=data.drop(columns=["Sex"],axis=1)
x.head()
```

```
Out[ ]:
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	12
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	4
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	6
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	7
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	4

9. Scale the independent variables

```
In [ ]: from sklearn.preprocessing import scale
X_Scaled = pd.DataFrame(scale(x), columns=x.columns)
X_Scaled.head()
```

```
Out[ ]:
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	-0.574558	-0.432149	-1.064424	-0.641898	-0.607685	-0.726212	-0.638217	1.555152
1	-1.448986	-1.439929	-1.183978	-1.230277	-1.170910	-1.205221	-1.212987	-0.884841
2	0.050033	0.122130	-0.107991	-0.309469	-0.463500	-0.356690	-0.207139	-0.274842
3	-0.699476	-0.432149	-0.347099	-0.637819	-0.648238	-0.607600	-0.602294	0.030157
4	-1.615544	-1.540707	-1.423087	-1.272086	-1.215968	-1.287337	-1.320757	-0.884841

10. Split the data into training and testing

```
In [ ]: from sklearn.model_selection import train_test_split
X_Train, X_Test, Y_Train, Y_Test = train_test_split(X_Scaled, y, test_size=0.2, random_state=0)
```

```
In [ ]: X_Train.shape, X_Test.shape
```

```
Out [ ]: ((3341, 8), (836, 8))
```

```
In [ ]: Y_Train.shape, Y_Test.shape
```

```
Out [ ]: ((3341,), (836,))
```

```
In [ ]: X_Train.head()
```

```
Out [ ]:
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
3141	-2.864726	-2.750043	-1.423087	-1.622870	-1.553902	-1.583867	-1.644065	-1.799838
3521	-2.573250	-2.598876	-2.020857	-1.606554	-1.551650	-1.565619	-1.626104	-1.494839
883	1.132658	1.230689	0.728888	1.145672	1.041436	0.286552	1.538726	1.555152
3627	1.590691	1.180300	1.446213	2.164373	2.661269	2.330326	1.377072	0.030157
2106	0.591345	0.474853	0.370226	0.432887	0.255175	0.272866	0.906479	1.250153

```
In [ ]: X_Test.head()
```

```
Out [ ]:
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
668	0.216591	0.172519	0.370226	0.181016	-0.368878	0.569396	0.690940	0.945154
1580	-0.199803	-0.079426	-0.466653	-0.433875	-0.443224	-0.343004	-0.325685	-0.579842
3784	0.799543	0.726798	0.370226	0.870348	0.755318	1.764639	0.565209	0.335156
463	-2.531611	-2.447709	-2.020857	-1.579022	-1.522362	-1.538247	-1.572219	-1.799838
2615	1.007740	0.928354	0.848442	1.390405	1.415417	1.778325	0.996287	0.640155

```
In [ ]: Y_Train.head()
```

```
Out [ ]:
```

3141	1
3521	1
883	2
3627	2
2106	2

Name: Sex, dtype: int64

```
In [ ]: Y_Test.head()
```

```
Out [ ]:
```

668	2
1580	1
3784	2
463	1
2615	2

Name: Sex, dtype: int64

11. Build the Model

```
In [ ]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(n_estimators=10, criterion='entropy')
```

```
In [ ]: model.fit(X_Train, Y_Train)
```

```
Out [ ]: RandomForestClassifier(criterion='entropy', n_estimators=10)
```

```
In [ ]: y_predict = model.predict(X_Test)
```

```
In [ ]: y_predict_train = model.predict(X_Train)
```

12. Train the Model

```
In [ ]: from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [ ]: print('Training accuracy: ', accuracy_score(Y_Train, y_predict_train))
```

```
Training accuracy: 0.9787488775815624
```

13. Test the Model

```
In [ ]: print('Testing accuracy: ',accuracy_score(Y_Test,y_predict))  
Testing accuracy: 0.5526315789473685
```

14. Measure the performance using Metrics

```
In [ ]: pd.crosstab(Y_Test,y_predict)
```

```
Out[ ]: col_0  0   1   2  
Sex  
0    122  29  98  
1     37 217  37  
2     120  53 123
```

```
In [ ]: print(classification_report(Y_Test,y_predict))
```

	precision	recall	f1-score	support
0	0.44	0.49	0.46	249
1	0.73	0.75	0.74	291
2	0.48	0.42	0.44	296
accuracy			0.55	836
macro avg	0.55	0.55	0.55	836
weighted avg	0.55	0.55	0.55	836